

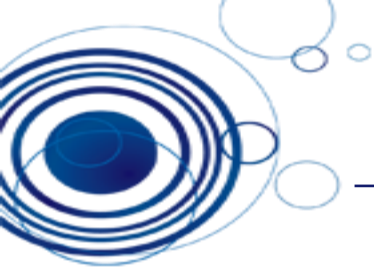
LOGO

BigData Engineering

10주차: Supervised Learning
강의 : 신경섭

11101001110000111110101110010101010011001010011010111101001110000111110101110010101010011001010011010111101001110
11100110000011011101001011101011111010101010010101111100110000011011101001011101011111010101010010101111100110000
110100111001101010101001011011111010100110101001010111010011100110101010100101101111010100110101001010111010011100

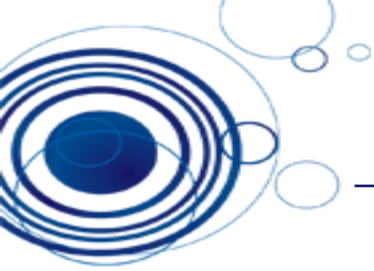




We covered...

- Data management by python
 - Numpy, pandas, data acquisition
- Machine learning workflow with data

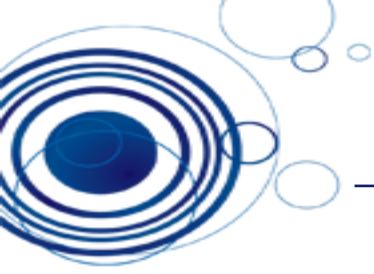




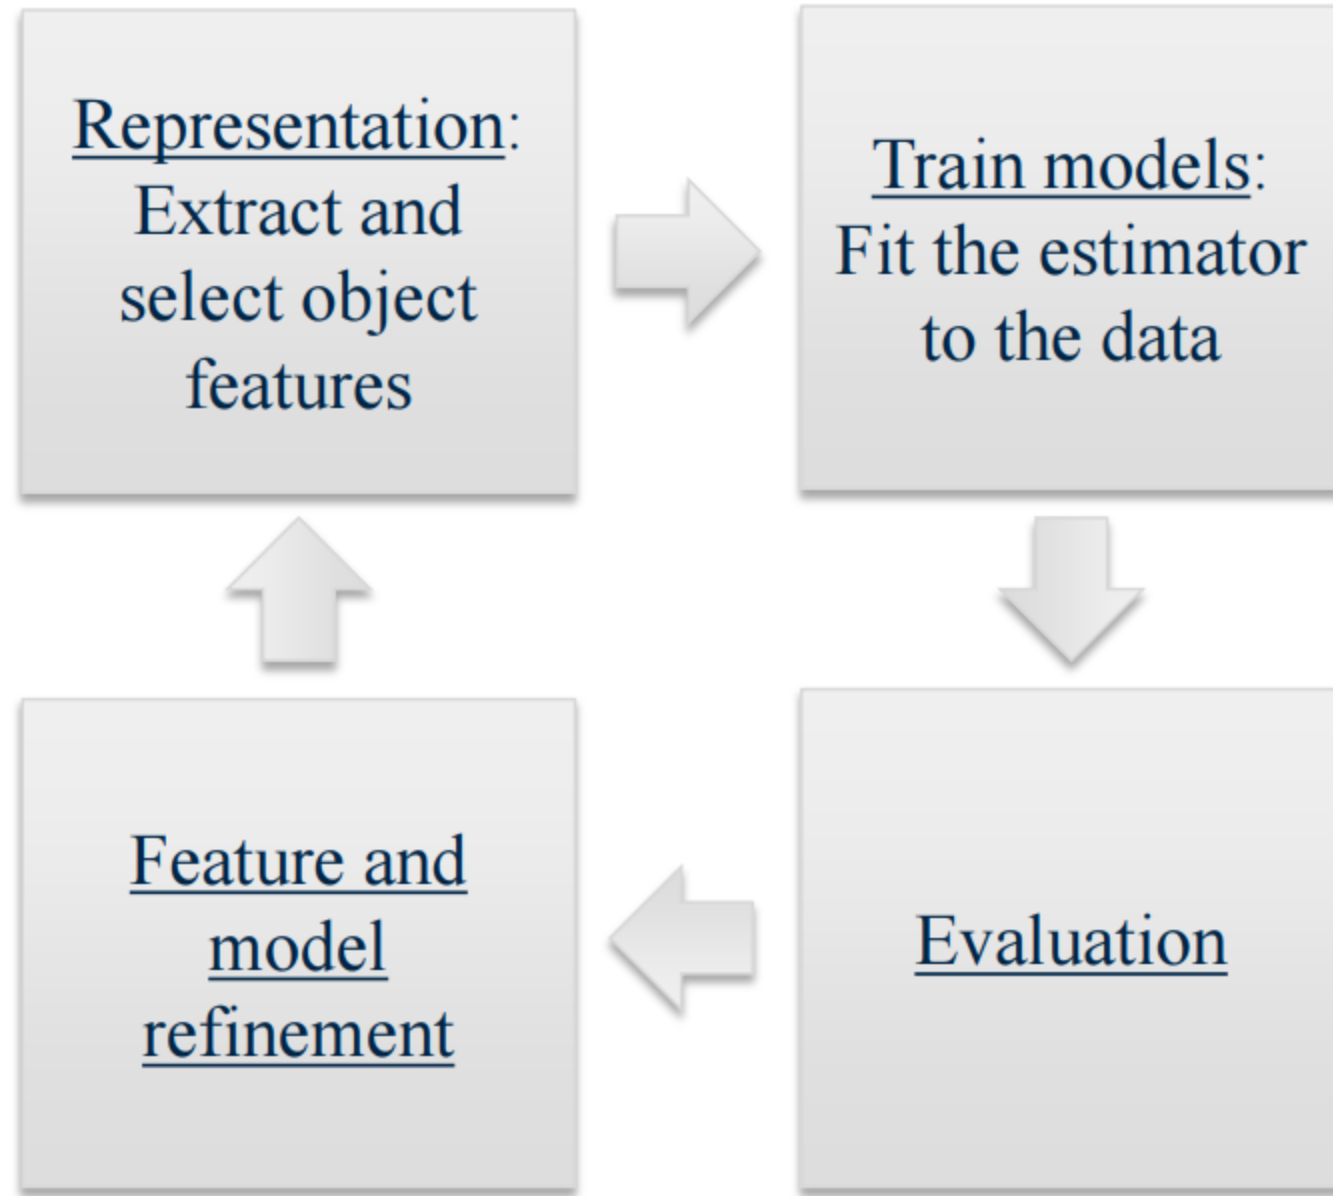
Today's Subjects

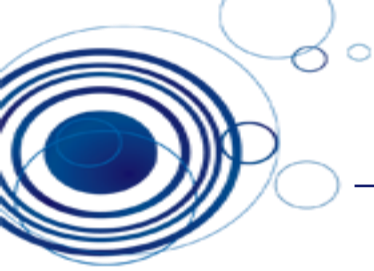
- EDA (Exploratory Data Analysis)
- Supervised learning
 - Examples – k-NN classifier
 - Regression : linear regression
 - Classification : logistic regression based binary classification





Represent / Train / Evaluate / Refine Cycle

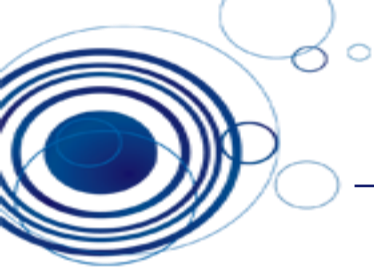




Before feature representation....

- EDA (Exploratory Data Analysis)
 - Understanding your variables
 - Cleaning your dataset
 - Analyzing relationship between variables
- Read: <https://towardsdatascience.com/an-extensive-guide-to-exploratory-data-analysis-ddd99a03199e>

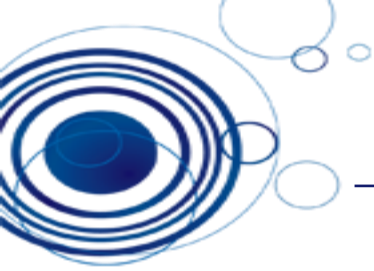




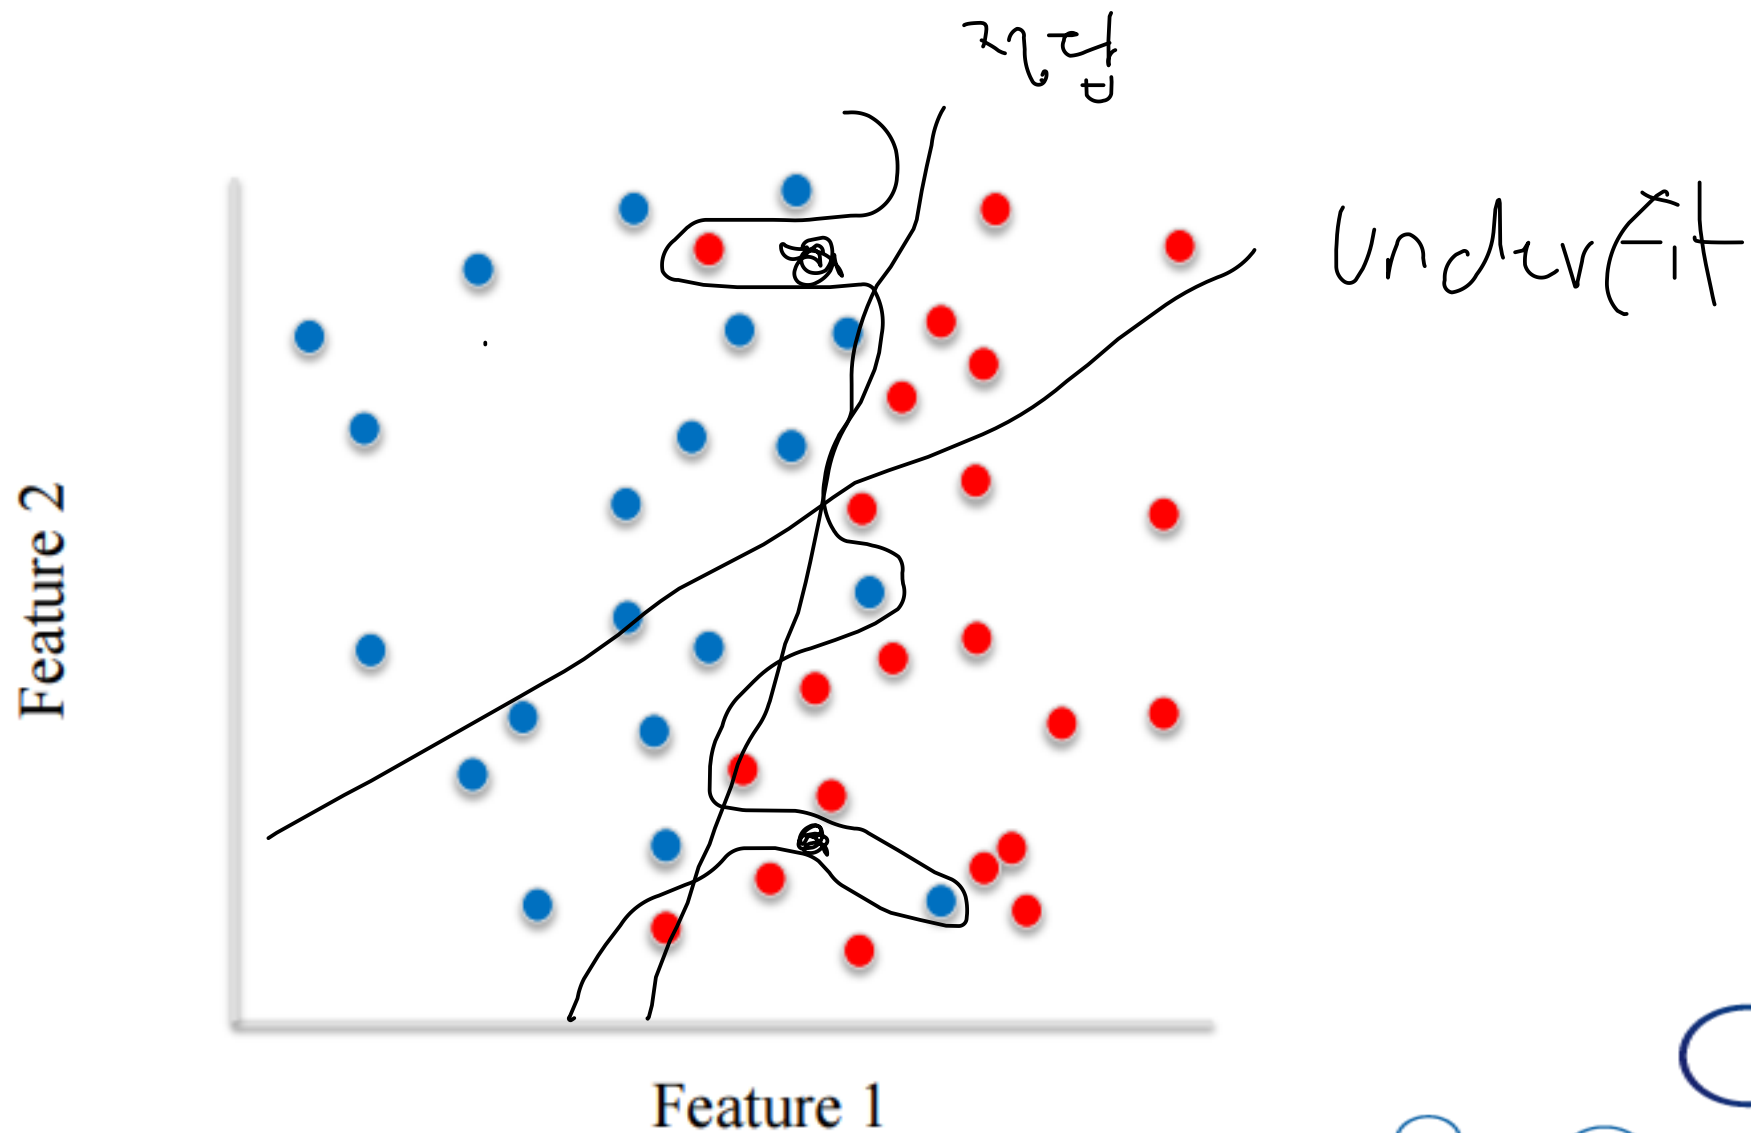
Generalization, overfitting, and underfitting


- Generalization ability refers to an algorithm's ability to give accurate predictions for new, previously unseen data.
- Assumptions:
 - Future unseen data (test set) will have the same properties as the current training sets.
 - Thus, models that are accurate on the training set are expected to be accurate on the test set.
 - But that may not happen if the trained model is tuned too specifically to the training set.
- Models that are too complex for the amount of training data available are said to overfit and are not likely to generalize well to new examples.
- Models that are too simple, that don't even do well on the training data, are said to underfit and also not likely to generalize well.






Overfitting in classification







Overfitting

- The relationship between model complexity and training/test performance





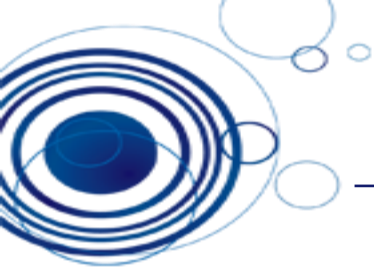
Fruit Dataset



	fruit_label	fruit_name	fruit_subtype	mass	width	height	color_score
0	1	apple	granny_smith	192	8.4	7.3	0.55
1	1	apple	granny_smith	180	8.0	6.8	0.59
2	1	apple	granny_smith	176	7.4	7.2	0.60
3	2	mandarin	mandarin	86	6.2	4.7	0.80
4	2	mandarin	mandarin	84	6.0	4.6	0.79
5	2	mandarin	mandarin	80	5.8	4.3	0.77
6	2	mandarin	mandarin	80	5.9	4.3	0.81
7	2	mandarin	mandarin	76	5.8	4.0	0.81
8	1	apple	braeburn	178	7.1	7.8	0.92
9	1	apple	braeburn	172	7.4	7.0	0.89
10	1	apple	braeburn	166	6.9	7.3	0.93
11	1	apple	braeburn	172	7.1	7.6	0.92
12	1	apple	braeburn	154	7.0	7.1	0.88
13	1	apple	golden_delicious	164	7.3	7.7	0.70
14	1	apple	golden_delicious	152	7.6	7.3	0.69
15	1	apple	golden_delicious	156	7.7	7.1	0.69
16	1	apple	golden_delicious	156	7.6	7.5	0.67


fruit_data_with_colors.txt






The k-Nearest Neighbor(k-NN) Algorithm


- Find k nearest neighbor data to classification



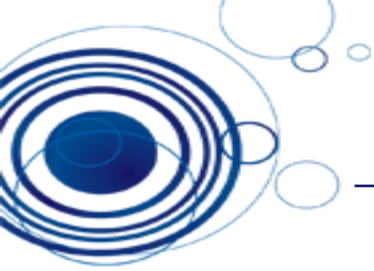
K=1



K=5



K=10




Linear regression

- Example : linear regression model with one variable

Input instance: $\mathbf{x} = (x_0)$

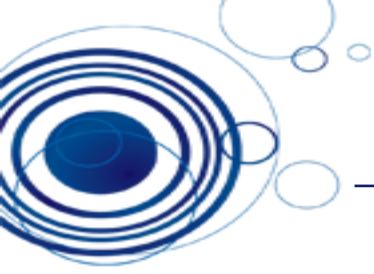
Predicted output: $\hat{y} = \widehat{w}_0 x_0 + \hat{b}$

Parameters to estimate: \widehat{w}_0 (slope)
 \hat{b} (y-intercept)

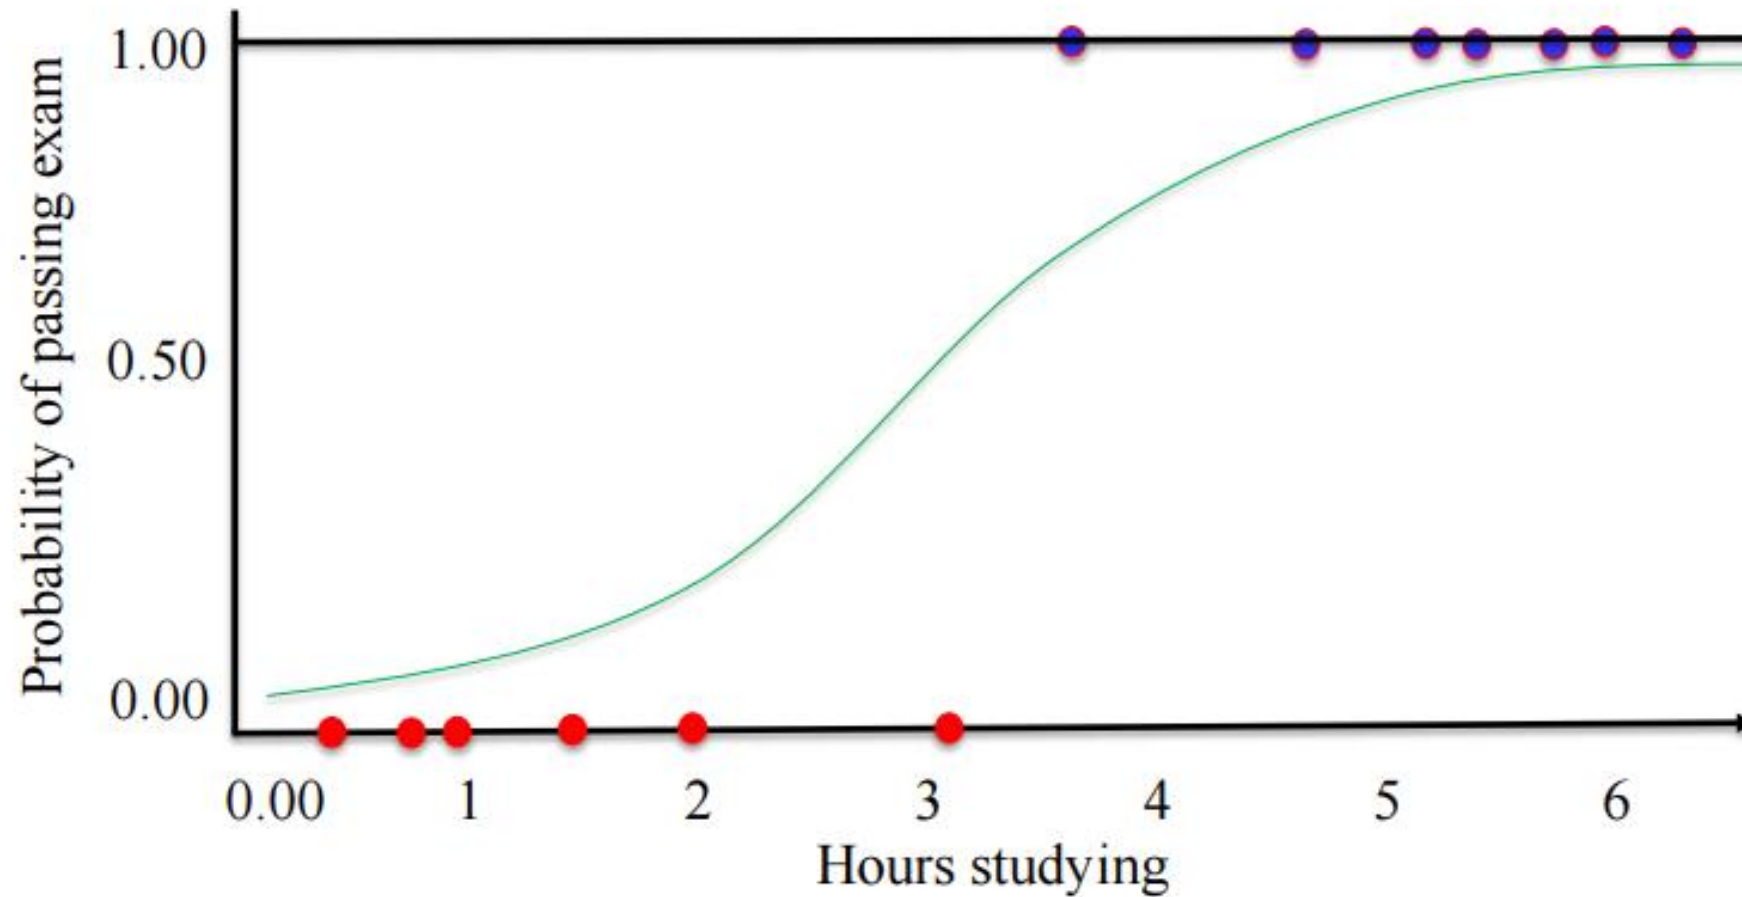


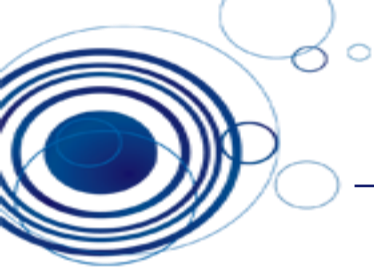
- Objective: minimize $RSS(\mathbf{w}, b) = \sum_{i=1}^N (y_i - (\mathbf{w} \cdot \mathbf{x}_i + b))^2$






Linear Regression to Logistic Regression





Linear models for classification:

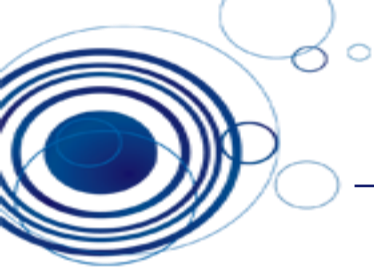
Logistic Regression



$$\hat{y} = \text{logistic}(\hat{b} + \hat{w}_1 \cdot x_1 + \cdots \hat{w}_n \cdot x_n)$$

$$= \frac{1}{1 + \exp[-(\hat{b} + \hat{w}_1 \cdot x_1 + \cdots \hat{w}_n \cdot x_n)]}$$





Multi-class classification with linear models

```
clf = LinearSVC(C=5, random_state = 67)
clf.fit(X_train, y_train)

print(clf.coef_)

[[-0.23401135  0.72246132]
 [-1.63231901  1.15222281]
 [ 0.0849835   0.31186707]
 [ 1.26189663 -1.68097    ]]

print(clf.intercept_)
[-3.31753728  1.19645936 -2.7468353  1.16107418]
```

