

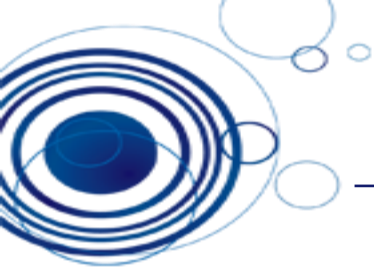
LOGO

BigData Engineering

12주차: Model Evaluation
강의 : 신경섭

11101001110000111110101110010101010011001010011010111101001110000111110101110010101010011001010011010111101001110
11100110000011011101001011101011111010101010010101111100110000011011101001011101011111010101010010101111100110000
110100111001101010101001011011111010100110101001010111010011100110101010100101101111010100110101001010111010011100

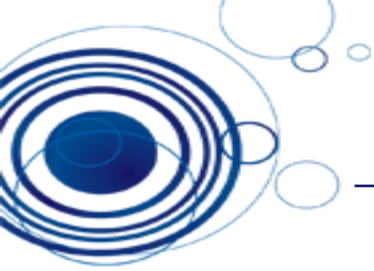




We covered...

- Data management/Visualization by python
 - Numpy, pandas, data acquisition
- Machine learning workflow with data
- EDA (Exploratory Data Analysis)
- Supervised learning
 - k-NN classifier
 - logistic regression based binary classification
 - Support vector machine
 - Decision tree
 - Random Forest

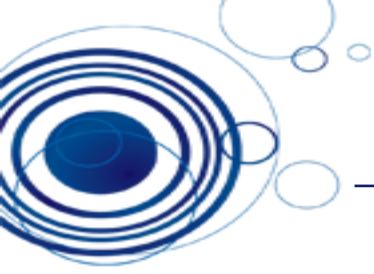




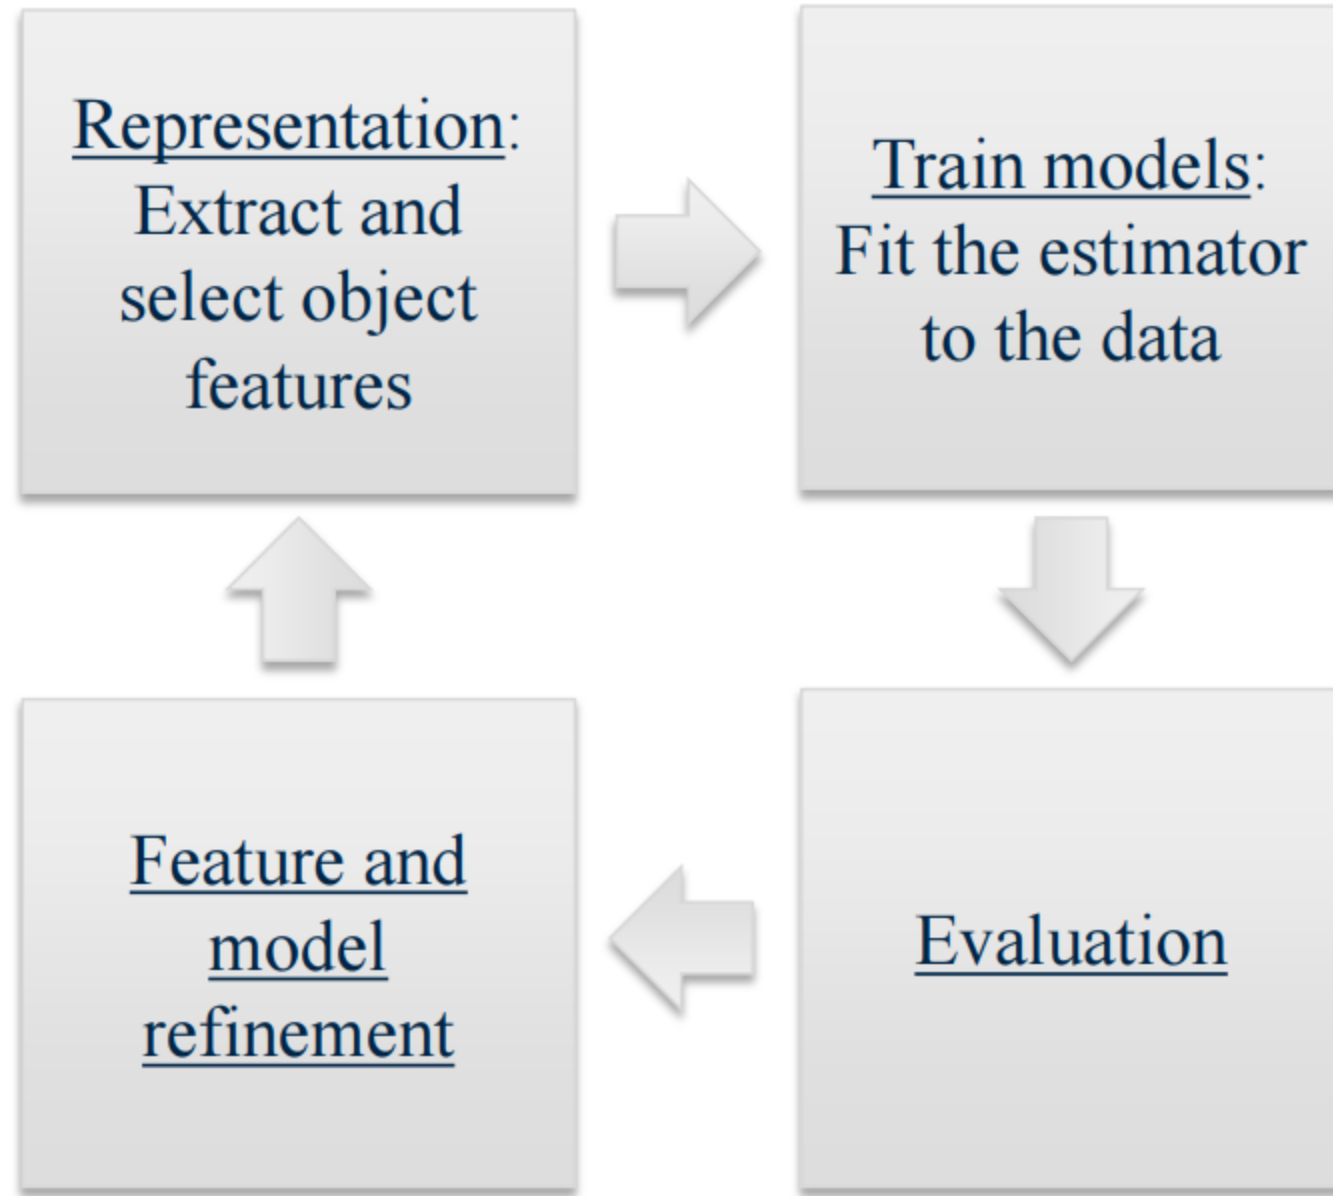
Today's Subjects

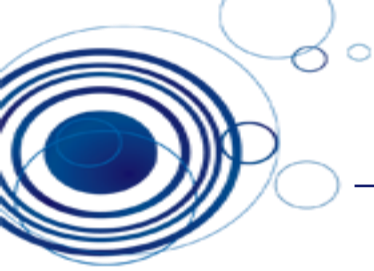
- Model Evaluation
 - Accuracy
 - Precision
 - Recall





Represent / Train / Evaluate / Refine Cycle

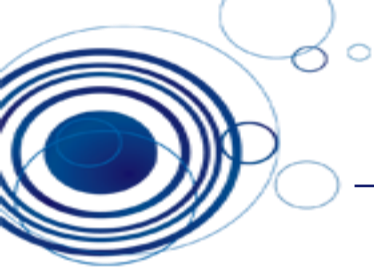




Evaluation

- Different applications have very different goals
- Accuracy is widely used, but many others are possible, e.g.:
 - User satisfaction (Web search)
 - Amount of revenue (e-commerce)
 - Increase in patient survival rates (medical)

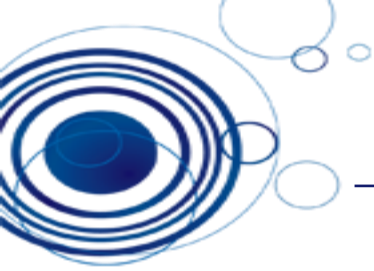




Evaluation

- It's very important to choose evaluation methods that match the goal of your application.
- Compute your selected evaluation metric for multiple different models.
- Then select the model with 'best' value of evaluation metric.






Accuracy with imbalanced classes

- Suppose you have two classes:
 - Relevant (R): the positive class
 - Not_Relevant (N): the negative class
- Out of 1000 randomly selected items, on average
 - One item is relevant and has an R label
 - The rest of the items (999 of them) are not relevant and labelled N.
- Recall that:

$$\text{Accuracy} = \frac{\text{\#correct predictions}}{\text{\#total instances}}$$

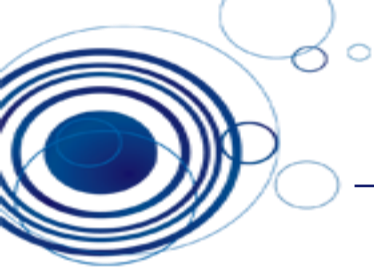




Accuracy with imbalanced classes

- You build a classifier to predict relevant items, and see that its accuracy on a test set is 99.9%.
 - Wow! Amazingly good, right?
- For comparison, suppose we had a "dummy" classifier that didn't look at the features at all, and always just blindly predicted the most frequent class (i.e. the negative N class).
- Assuming a test set of 1000 instances, what would this dummy classifier's accuracy be?

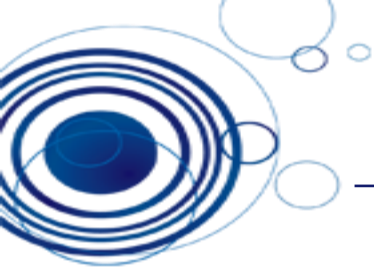




Dummy classifiers

- Some commonly-used settings for the strategy parameter for DummyClassifier in scikit-learn:
 - most_frequent : predicts the most frequent label in the training set.
 - stratified : random predictions based on training set class distribution.
 - uniform : generates predictions uniformly at random.
 - constant : always predicts a constant label provided by the user.





Binary prediction outcomes

<u>True</u> negative	TN	FP
<u>True</u> positive	FN	TP
	<u>Predicted</u> negative	<u>Predicted</u> positive

Label 1 = positive class
(class of interest)

Label 0 = negative class
(everything else)

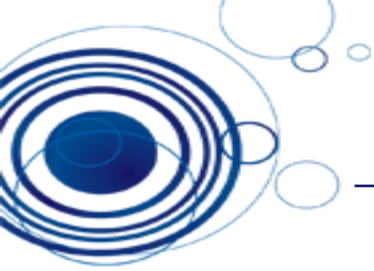
TP = true positive

FP = false positive (Type I error)

TN = true negative

FN = false negative (Type II error)



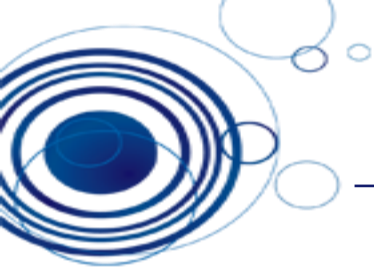


Confusion Matrix for Binary Prediction Task


True negative	TN = 400	FP = 7	
True positive	FN = 17	TP = 26	
	Predicted negative	Predicted positive	$N = 450$

Always look at the confusion matrix for your classifier.



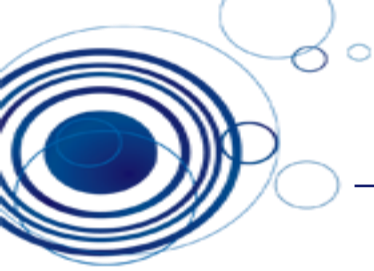


Visualization of Different Error Types



$TN = 429$	$FP = 6$
$FN = 2$	$TP = 13$





Model Evaluation Metric: Accuracy

True
negative

TN = 400

FP = 7

$$\text{Accuracy} = \frac{TN+TP}{TN+TP+FN+FP}$$

$$= \frac{400+26}{400+26+17+7}$$

$$= 0.95$$

True
positive

FN = 17


TP = 26

$N = 450$

Predicted
negative

Predicted
positive





Model Evaluation Metric: Classification Error

True
negative

TN = 400

FP = 7

True
positive

FN = 17

TP = 26

Predicted
negative

Predicted
positive


$N = 450$

$$\text{ClassificationError} = \frac{FP + FN}{TN + TP + FN + FP}$$

$$= \frac{7+17}{400+26+17+7}$$

$$= 0.060$$





Model Evaluation Metric: Precision


True negative	TN = 400	FP = 7	
True positive	FN = 17	TP = 26	
	Predicted negative	Predicted positive	$N = 450$

$$\text{Precision} = \frac{TP}{TP+FP}$$

$$= \frac{26}{26+7}$$

$$= 0.79$$





Model Evaluation Metric: Recall

True
negative

TN = 400

FP = 7

True
positive

FN = 17

TP = 26

Predicted
negative

Predicted
positive

$N = 450$

$$\text{Recall} = \frac{TP}{TP+FN}$$


$$= \frac{26}{26+17}$$

$$= 0.60$$

Recall is also known as:

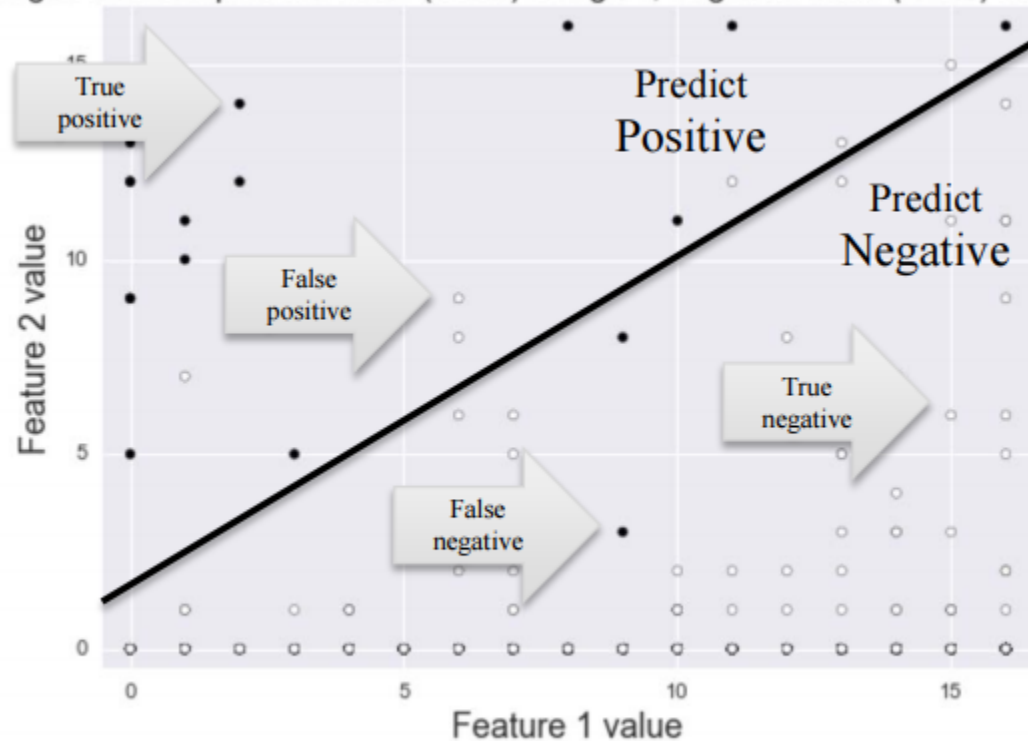
- True Positive Rate (TPR)
- Sensitivity
- Probability of detection





The Precision-Recall tradeoff

digits dataset: positive class (black) is digit 1, negative class (white) all others




TN = 429	FP = 6
FN = 2	TP = 13

$$\text{Precision} = \frac{TP}{TP+FP} = \frac{13}{19} = 0.68$$


$$\text{Recall} = \frac{TP}{TP+FN} = \frac{13}{15} = 0.87$$





The Precision-Recall tradeoff

digits dataset: positive class (black) is digit 1, negative class (white) all others




TN = 435	FP = 0
FN = 8	TP = 7

$$\text{Precision} = \frac{TP}{TP+FP} = \frac{7}{7} = 1.00$$


$$\text{Recall} = \frac{TP}{TP+FN} = \frac{7}{15} = 0.47$$





The Precision-Recall tradeoff

digits dataset: positive class (black) is digit 1, negative class (white) all others




TN = 408	FP = 27
FN = 0	TP = 15

$$\text{Precision} = \frac{TP}{TP+FP} = \frac{15}{42} = 0.36$$

$$\text{Recall} = \frac{TP}{TP+FN} = \frac{15}{15} = 1.00$$





The Precision-Recall tradeoff

- Recall-oriented machine learning tasks:
 - Search and information extraction in legal discovery
 - Tumor detection
 - Often paired with a human expert to filter out false positives
- Precision-oriented machine learning tasks:
 - Search engine ranking, query suggestion
 - Document classification
 - Many customer-facing tasks (users remember failures!)

