

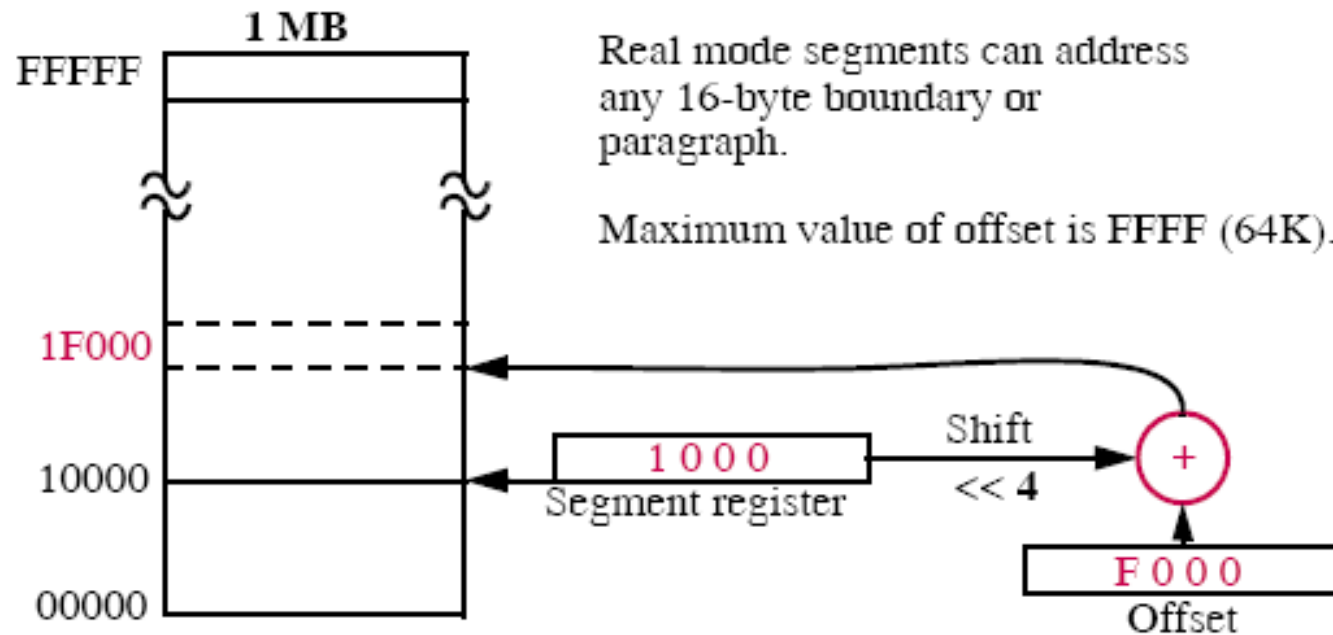
# 12 Linux User Programs

# Multitasking

- Multitasking
  - Allows programs to take turns using processor
  - Scheduler decides to force user programs to yield to one another even when they are not finished
  - Task state segment: bumped process saves its current state into a region in memory
- Memory protection
  - Multitasking requires memory protection
  - Linux uses paging to implement a memory protection
  - Every memory reference in a user program is turned into a request for memory access which is required to go across the desk of the paging system for approval

# Cf) Real Mode Memory Addressing

- Only mode available to the 8086 and 8088.
  - Allow the processor to address only the first 1MB of memory.
  - DOS requires real mode.
- *Segments and Offsets:*
  - Effective address = Segment address + an offset.



# Protected Mode

- To change to protected mode from DOS

```
MOV EAX, CR0  
OR EAX, 1      : Make the least significant bit 1.  
MOV CR0, EAX   : Here goes!
```

- x86 CR0 Register

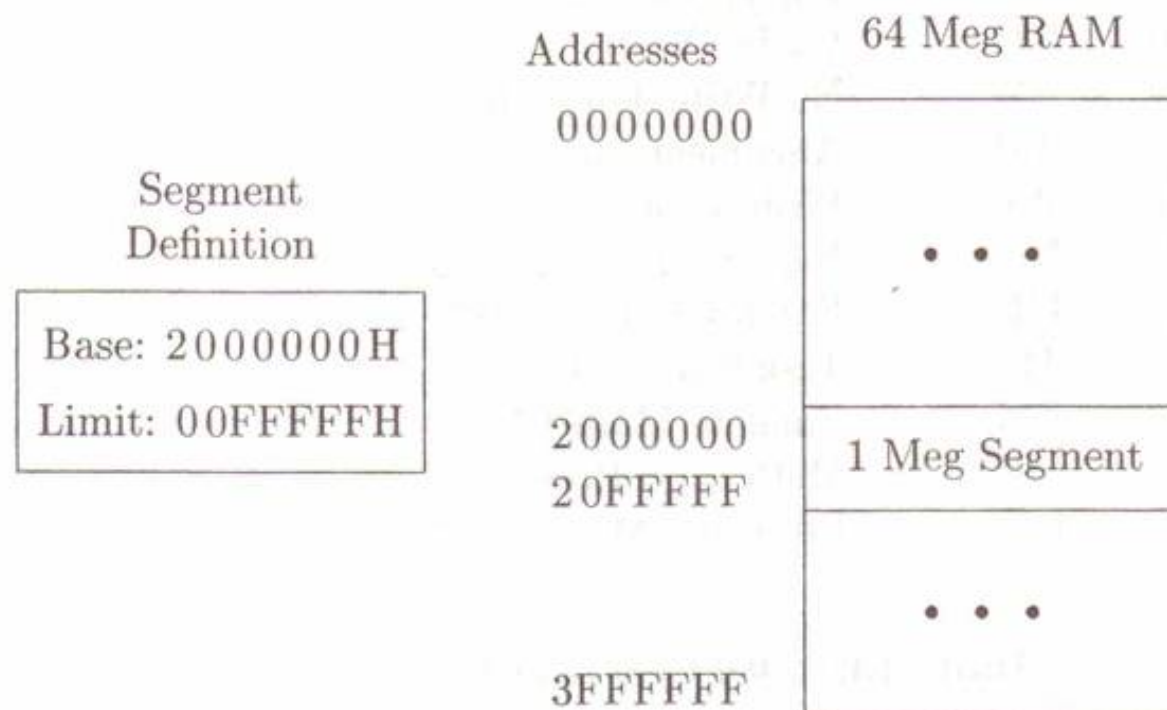
Bits	Label	Full Name
31	PG	Paging Enable
30	CD	Cache Disable
29	NW	No Write Through
18	AM	Alignment Mask
15	WP	Write Protect
5	NE	Numeric Error Enable
4	ET	Extension Type (287 vs. 387)
3	TS	Task Switched
2	EM	Emulate Math Chip
1	MP	Math Chip Present
0	PE	Protected Mode Enable

# Protected Mode

- In protected mode
  - all memory references depend on *global descriptor table*
  - instructions are fetched from a different segment
  - each entry in the interrupt table is a specially formatted eight-byte entry; before going into protected mode, a new interrupt table must be created

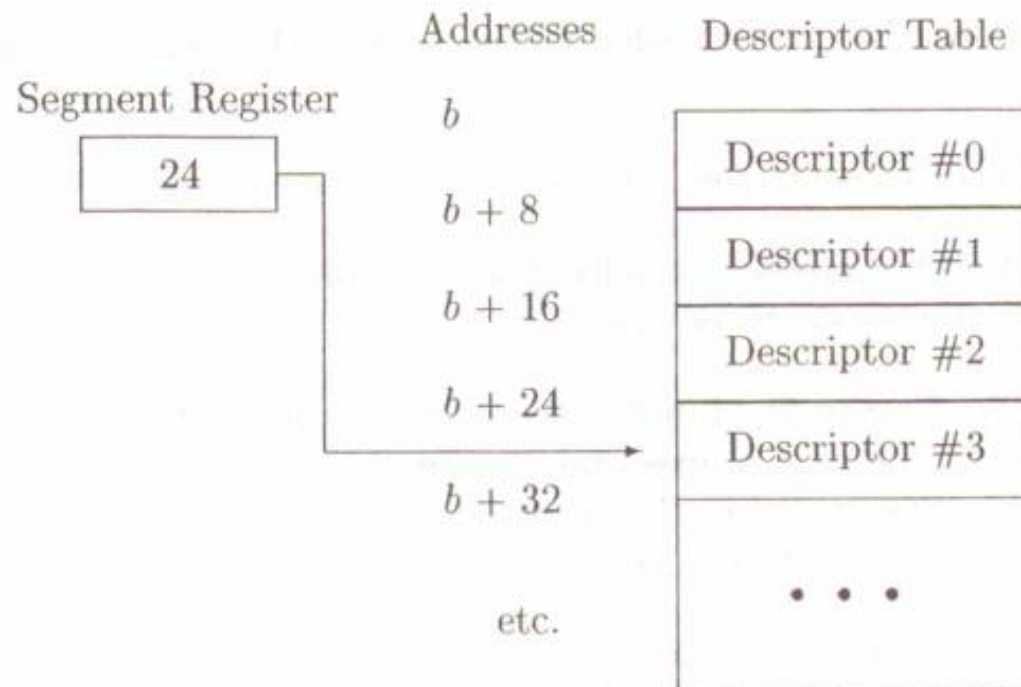
# Protected Mode Segmentation

- A protected mode memory segment



# Protected Mode Segmentation

- segment, segment descriptor, descriptor table, selector



# Protected Mode Memory Segments

- Characteristics of a protected mode memory segment
  - Base : the address of the first byte of the segment
  - Limit: the last valid offset address
  - Access permissions: permissions to read, write, and execute can be specified
  - Access privilege level: the least privileged value that CPL (current privilege level) may have in order for access to be granted



# Protected Mode Memory Segments

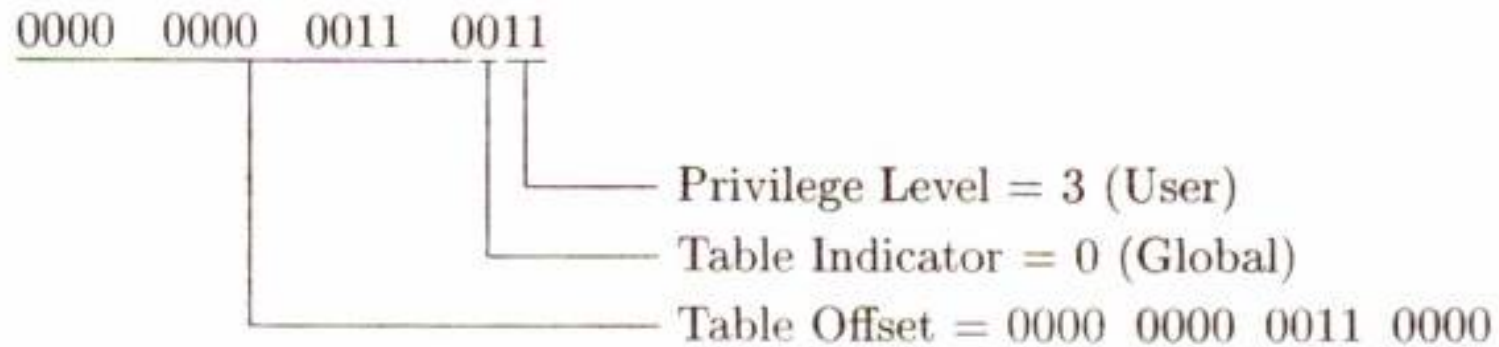
- Selector, base, and limit registers

	Segment Register	Base Register	Limit Register
ES			
CS			
SS			
DS			
FS			
GS			
LDTs			
TSS			
GDT			
IDT			

- Special memory segments, GDT and IDT

# Selectors

- Selector encoding



# Segment Operations

- Segment load and store commands

	Load	Store	MOV Load	MOV Store
ES	LES		X	X
CS				X
SS	LSS		X	X
DS	LDS		X	X
FS	LFS		X	X
GS	LGS		X	X
LDTs	LLDT	SLDT		
TSS	LTR	STR		
GDT	LGDT	SGDT		
IDT	LIDT	SIDT		

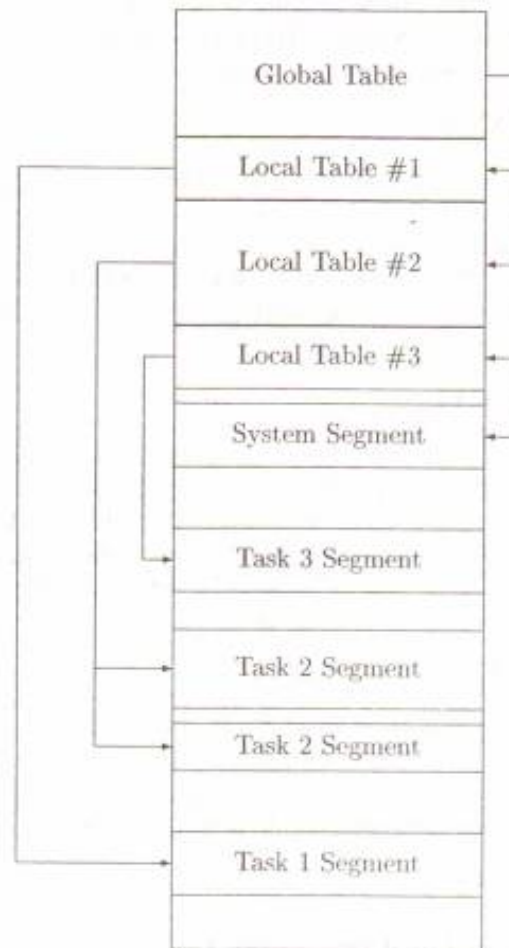
# Descriptor Encoding

- Segment descriptor layout

63 56 55 54				51 48 47			44 43			40 39			16 15		0	
Top Base Byte				G	D	Top Limit Bits		P	D P L	1	E	R W	A	Base Address Lower 3 Bytes		Limit Lower 2 Bytes

# Task Isolation with LDT

- Memory with protected segments



# Task Isolation with LDT

- Consider this code

```
MOV AX, 0  
; AX is a 16-bit selector.  
; Bit 2 is 0. ; We're pointing into the global table.  
; Bits 0 and 1 are 0  
; That's system level privilege.  
HAK: LLDT AX  
; AX is pointing at some descriptor  
; It could be a local table descriptor  
; If so it's probably not our own local table.  
; And in that case we have just borrowed someone  
; else's private descriptor table and we can  
; access anything listed in there.  
ADD AX, 8  
; Maybe not.  
; Go to the next descriptor  
JMP HAK  
; The local tables are in here.  
; We'll find them.
```

- Memory protection by itself doesn't work; user processes must not be allowed to execute code like the above

# Setting Up the GDT

```
MOV AX, 2000H ; Set up a segment for the scratch area.
MOV DS, AX    ;
MOV AX, 4FH   ; GDT's limit is a 2 byte value.
MOV [DS:0], AX ; Store 79 into the first 2 bytes of
               ; scratch.
MOV EAX, 10000H ; GDT's base address is a 4 byte value.
MOV [DS:2], EAX ; Store this 4 byte address
               ; 2 bytes into the scratch.
               ; (Don't overwrite the first two bytes.)
LGDT [DS:0]    ;
; This loads the both the GDT base address register
               ; and the GDT limit register
               ; i.e., load all 6 bytes.
```

# Paging

- x86 I/O address vs. memory address

IN AL,[DX] ; where DX=FC30H

→ 16 bits of DX is directly applied to address bus

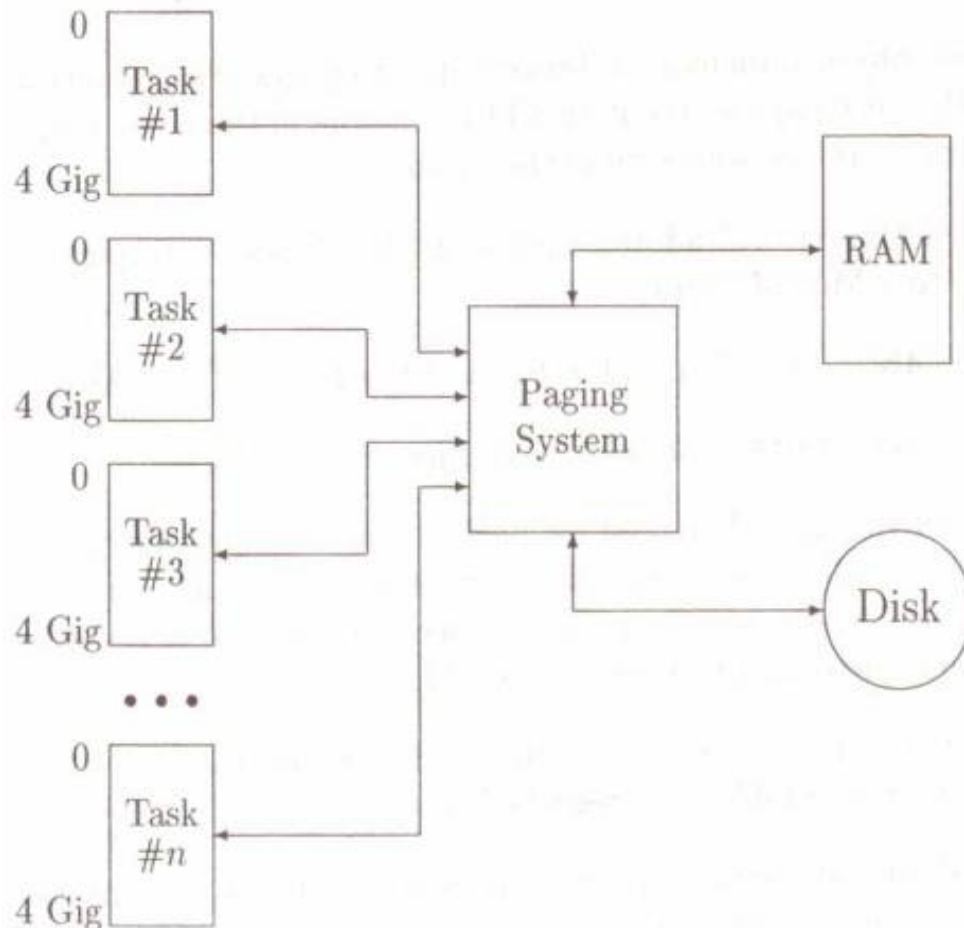
MOV AL,[0FC30H]

→ memory address FC30 undergo a two-step process of *segmentation* and *paging* \*

\* Linux uses only paging



# Virtual Memory



- x86 virtual memory system
  - Each process has address space of 0 ~ 4G
  - Process's address space is broken up into *pages*
  - Only parts of pages are loaded into memory
  - Each process has its own set of *page tables* to translate the virtual addresses into physical ones

# 4K pages


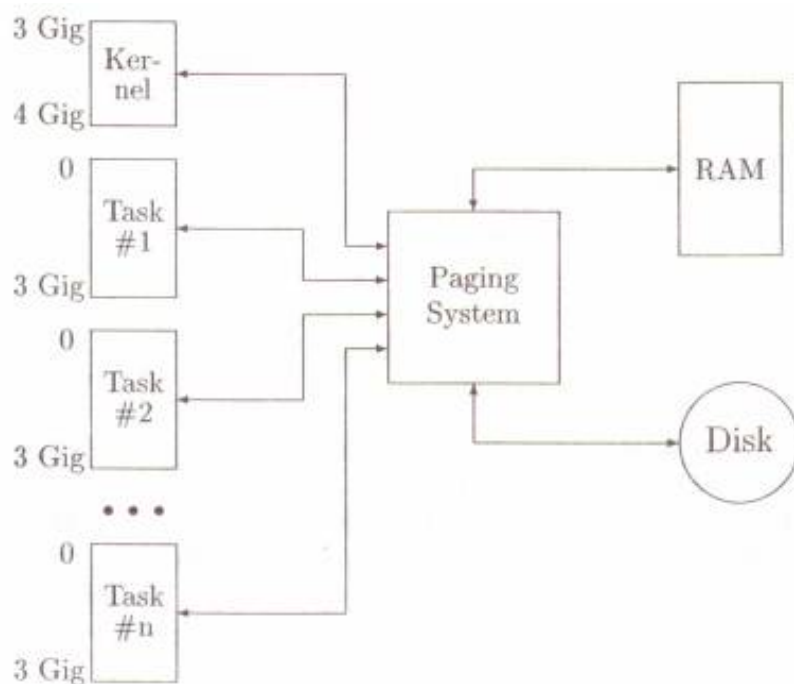
- In order to set up a paging system
  - The available physical memory is broken up into 4K size page frames
  - Each task's 4 Gig virtual address space is also divided into pages
  - The portion of the disk used for paging, called swap space, is also divided into pages
  - The page tables used to translate virtual addresses are set up in units of pages

# Paging vs. swapping

- Swapping
  - Copying entire process to and from the disk
- Page faults
  - When a machine instruction refers to a memory location actually on disk, this is called *page fault*
  - When a page fault occurs, some page frame in RAM has to be chosen so that the needed page may be loaded into it from the disk
  - If no page frame is free, then some page of memory has to be chosen for transfer to the disk so as to free up a page frame


# Kernel addresses

- Linux virtual memory system
- Single process view of memory



# Address Translation

- Two-tier page tables



# Kernel


- Kernel pages
  - The Linux paging never pages kernel memory out to disk
  - The Linux kernel is therefore always in RAM
- Kernel modules
  - It has been possible to add and delete kernel memory
  - *insmod* command install module
  - *rmmod* command remove module
  - *free* inspect memory to find out the amount of memory available

# Address Translation

- x86 paging obtains its translation from page tables set up by the operating system
- Since the translation process is hard-wired into the processor, x86 page tables must have the two-tiered structure
- X86 page table
  - Each table has 1024 four-byte entries -> 4K size -> same to page size
  - page table lists the location of 1024 page tables
  - Total memory referenced by a page table is  $1024 * 1024 * 4096 = 4 \text{ GB}$  -> size of virtual address space

# Address Translation

- Address parsing




- Example of a virtual address of '1A2B3C4DH'

1	A	2	B	3	C	4	D	
0001	1010	0010	1011	0011	1100	0100	1101	virtual address
0001	1010	00						index into page table table
		10	1011	0011				index into page table
				1100	0100	1101		index into page

- 104<sup>th</sup> entry of the page table table
- 691<sup>st</sup> entry of the page table
- Offset is 0C4DH
- If the beginning of the page is 12340000H the effective address is 1234C4DH



# Page table entries



- PP : page is located in RAM
- If PP is not set
  - It is located on disk -> page fault
  - The address is not valid at all -> segmentation fault
- CR3 register
  - The physical address of the page table table

# Program Segments

- Non-writeable memory
  - W\_R bit
    - Set means that page is writeable
    - Clear means not writeable
  - Attempt to write to not writeable memory -> segmentation fault
- Writeable memory
  - At least two different pages are needed

# Program Segments

- Segments (in executable file)
  - Text segment (W\_R bit is cleared)
  - Data segment (W\_R bit is set)
- Sections (in object file)
  - An executable file has segments, an object code file has sections
- Section definition directives
  - `section .text`
  - `section .data`

# Program with section

- Program produce .o file with both text and data section

```
; Tiny program which increments a byte of stored data.  
;  
        global main  
        section .text  
main:    MOV AL,[XYZ]  
        INC AL  
        MOV [XYZ], AL  
        RET  
        section .data  
XYZ:     db 3
```

- Segments in C code

```
char a = 'A';
main()
{   char *n;
    int i,t;

    n = &a;
    printf("Enter a pointer displacement in hex:");
    scanf("%x",&i);
    printf("Pointer value in hex = %x\n", (n + i));
    printf("Read, write, or skip?  (Enter 0, 1, or 2):");
    scanf("%d", &t);
    if (t == 0)
        printf("Contents there= %x", *(n + i));
    else if (t == 1)
        *(n + i) = 'Y';
    printf("\n");
}
```

- Calling C from Assembler

```
; This program calculates y = 2 * x + 1
    section .text
    global main
    extern printf
    extern scanf
main:  PUSH ABC
      CALL printf
      ADD ESP, 4
      PUSH XYZ
      PUSH BCD
      CALL scanf
      ADD ESP, 8
      MOV EAX, [XYZ]
      ADD EAX, EAX
      INC EAX
      PUSH EAX
      PUSH CDE
      CALL printf
      ADD ESP, 8
      RET
ABC:  db "Enter a number: ",0AH,0
BCD:  db "%d \x0",0AH,0
CDE:  db "You get %d.\x0",0AH,0
      section .data
XYZ:  db 4 * 0
```

# Other Data Segments

- Initialized data space

```
section .data
```

- Uninitialized data space

```
int x[10000];
```

- Dynamically allocated memory

```
malloc()
```

# ELF Format

ELF Header
Segment 1
Segment 2
Segment 3
Segment 4
Program Header Table
Section Header Table
Section 1
• • •
Section n

- Executable files in ELF format
- ELF header
  - Identifier string
  - File type
  - Machine architecture
  - Size, location, ...
- Program header
  - Type
  - File offset
  - Virtual address
  - Physical address
  - File size
  - Memory size
  - Permissions
  - Alignment



# ELF Format

- Program header table for chmod

0034:00d4 Program Header Table

	Seg	File	Virtual	Physical	File	Mem	Align
Num	Type	Offst	Address	Address	Size	Size	Perm -ment
00:	0006	0034	8048034	8048034	00a0	00a0	0005 0004
01:	0003	00d4	80480d4	80480d4	0013	0013	0004 0001
02:	0001	0000	8048000	8048000	1d29	1d29	0005 1000
03:	0001	1d30	804ad30	804ad30	011c	01fc	0006 1000
04:	0002	1dc4	804adc4	804adc4	0088	0088	0006 0004

Type 1 = loadable, 2 = dynamic info, 3 = interpreter  
4 = note, 6 = program header


# ELF Format

- ELF segments

```
int a = 0x55555555, b = 0x66666666;  
int c = 0x77777777, d = 0x88888888;  
  
main()  
{  
  
    /* Store numbers which will be very easy to pick out  
    from the machine code of the compiled program.  */  
  
    a = 0x11111111;  
    b = 0x22222222;  
    c = 0x33333333;  
    d = 0x44444444;  
  
    printf("Variables:  Addresses:\n")  
    printf(" a %x\n", &a);  
    printf(" b %x\n", &b);  
    printf(" c %x\n", &c);  
    printf(" d %x\n", &d);  
  
    /* %x means to use "hexadecimal" output.  */  
  
    printf("\n");  
}
```

# Object Files in ELF Format

- ELF object file format
  - Section header table
  - Section



- Section header table

01bc:0374 Section Header Table

	Name	Type	Flgs	Address	set	Size	Link	Info	Mod	Size	Align	Entry
00:	0000	0000	0000	00000000	0000	0000	0000	0000	0000	0000	0000	0000
01:	001b	0001	0006	00000000	0034	0091	0000	0000	0004	0000	0000	0000
02:	0021	0009	0000	00000000	0490	00a0	0009	0001	0004	0008	0000	0000
03:	002b	0001	0003	00000000	00c8	0010	0000	0000	0004	0000	0000	0000
04:	0031	0008	0003	00000000	00d8	0000	0000	0000	0004	0000	0000	0000
05:	0036	0007	0000	00000000	00d8	0014	0000	0000	0001	0000	0000	0000
06:	003c	0001	0002	00000000	00ec	0070	0000	0000	0001	0000	0000	0000
07:	0044	0001	0000	00000000	015c	0012	0000	0000	0001	0000	0000	0000
08:	0011	0003	0000	00000000	016e	004d	0000	0000	0001	0000	0000	0000
09:	0001	0002	0000	00000000	0374	00f0	000a	0009	0004	0010	0000	0000
0a:	0009	0003	0000	00000000	0464	002b	0000	0000	0001	0000	0000	0000

Name: index into the Section Header String Table (.shstrtab)

Type: 1 = Program Bits, 2 = Symbol Table, 3 = String Table

4 = Relocation Table, 5 = Hash Table, 6 = Dynamic Info

7 = Note, 8 = No Bits, 9 = Relocation Table