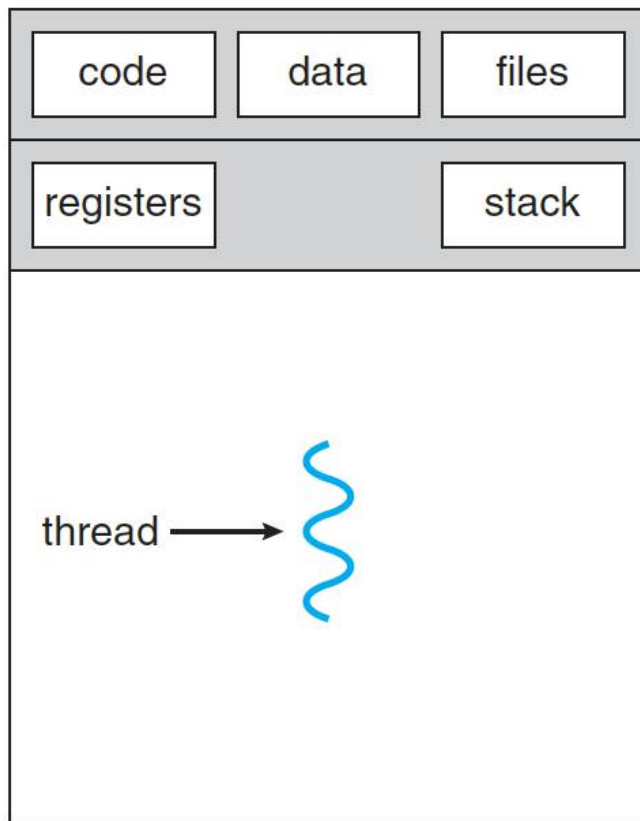


4 Threads

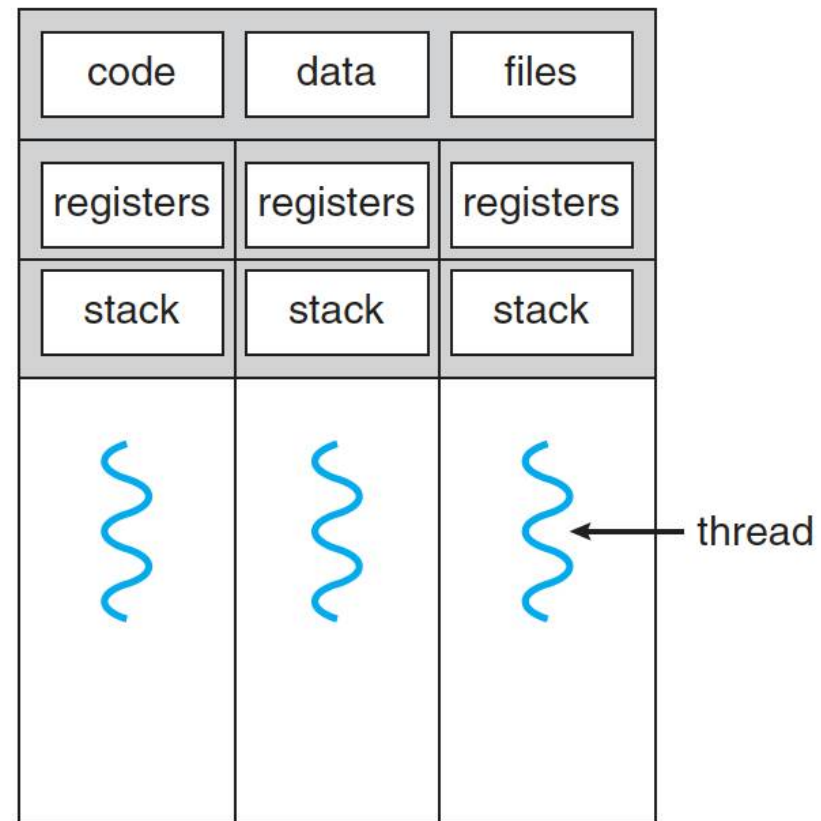
Thread

- A thread is a basic unit of CPU utilization
- It comprises a thread ID, a program counter, a register set and a stack
- A thread shares with other threads belonging to the same process its code section, data section, and other OS resources, such as open files and signals
- A traditional process has a single thread of control
- If a process has multiple threads of controls, it can perform more than one task at a time

Single-threaded vs. Multi-threaded Process



single-threaded process

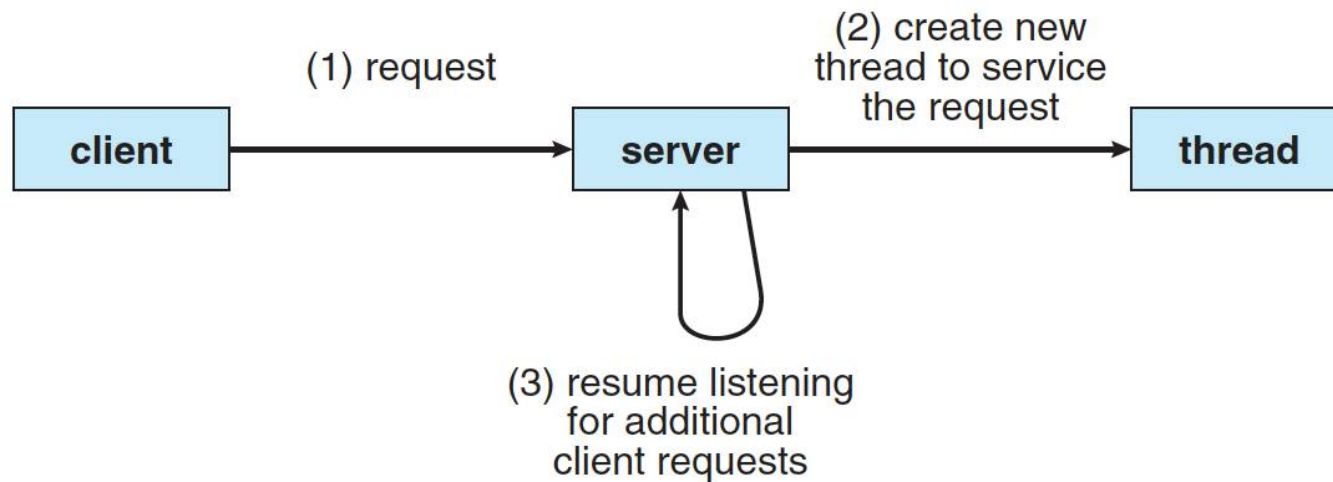


multithreaded process

Why Threads?

- A web browser might have one thread display images/text while another thread retrieves data from the network
- A word processor may have a thread for displaying graphics, another thread for responding to keystrokes from the user, and a third thread for performing spelling and grammar checking in the background
- Multithreaded web server architecture

Multithreaded Server Architecture



Benefits

- Responsiveness
 - Multithreading an interactive application may allow a program to continue running even if part of it is blocked or is performing a lengthy operation
- Resource Sharing
 - threads share the memory and the resources of the process to which they belong by default
- Economy
 - it is more economical to create and context-switch threads
- Scalability
 - threads may be running in parallel on different processing cores


Multithreading Models

- Support for threads may be provided at the user level, for *user threads*, or by the kernel, for *kernel threads*
- User threads are supported above the kernel and are managed without kernel support
- Kernel threads are supported and managed directly by the operating system
- Virtually all contemporary operating systems - including Windows, Linux, Mac OS X - support kernel threads
- A relationship must exist between user threads and kernel threads

Many-to-One

- Many user-level threads mapped to single kernel thread
 - Thread management is done by the thread library in user space, so it is efficient
 - However, the entire process will block if a thread makes a blocking system call.
 - Also, because only one thread can access the kernel at a time, multiple threads are unable to run in parallel on multicore systems
- Examples:
 - Solaris Green Threads
 - GNU Portable Threads


Many-to-One Model



One-to-One

- Each user-level thread maps to kernel thread
 - provides more concurrency than the many-to-one model by allowing another thread to run when a thread makes a blocking system call
 - It also allows multiple threads to run in parallel on multiprocessors
 - The only drawback to this model is that creating a user thread requires creating the corresponding kernel thread. Because the overhead of creating kernel threads can burden the performance of an application
- Examples
 - Windows NT/XP/2000
 - Linux
 - Solaris 9 and later

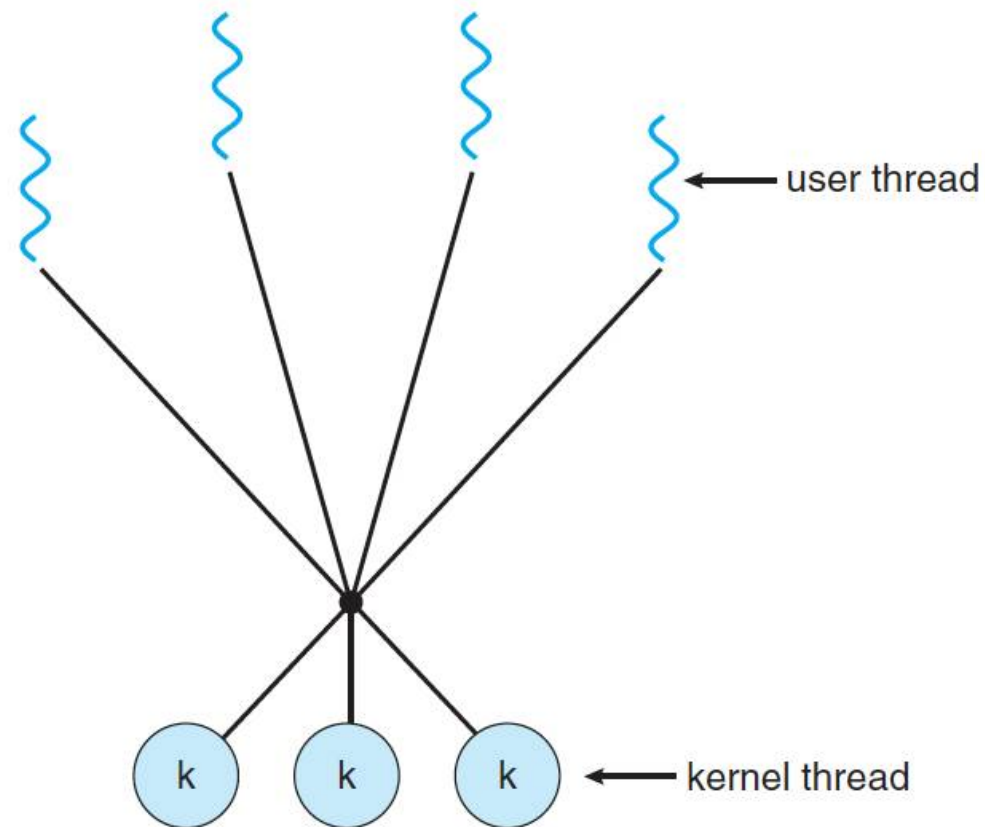
One-to-one Model




Many-to-Many Model

- The many-to-many model multiplexes many user-level threads to a smaller or equal number of kernel threads.
- Two-level model: Solaris prior to version 9
- Windows NT/2000 with the ThreadFiber package

Many-to-Many Model



Two-level Model



Thread Libraries

- A thread library provides the programmer with an API for creating and managing threads
- User-level library vs. kernel-level library
- Three main thread libraries are in use: POSIX Pthreads, Windows threads, Java threads

Pthreads

```
int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    if (argc != 2) {
        fprintf(stderr, "usage: a.out <integer value>\n");
        return -1;
    }
    if (atoi(argv[1]) < 0) {
        fprintf(stderr, "%d must be >= 0\n", atoi(argv[1]));
        return -1;
    }

    /* get the default attributes */
    pthread_attr_init(&attr);
    /* create the thread */
    pthread_create(&tid, &attr, runner, argv[1]);
    /* wait for the thread to exit */
    pthread_join(tid, NULL);

    printf("sum = %d\n", sum);
}

/* The thread will begin control in this function */
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; i++)
        sum += i;

    pthread_exit(0);
}
```


Windows Threads

```
#include <windows.h>
#include <stdio.h>
DWORD Sum; /* data is shared by the thread(s) */

/* the thread runs in this separate function */
DWORD WINAPI Summation(LPVOID Param)
{
    DWORD Upper = *(DWORD*)Param;
    for (DWORD i = 0; i <= Upper; i++)
        Sum += i;
    return 0;
}

int main(int argc, char *argv[])
{
    DWORD ThreadId;
    HANDLE ThreadHandle;
    int Param;

    if (argc != 2) {
        fprintf(stderr, "An integer parameter is required\n");
        return -1;
    }
    Param = atoi(argv[1]);
    if (Param < 0) {
        fprintf(stderr, "An integer >= 0 is required\n");
        return -1;
    }

    /* create the thread */
    ThreadHandle = CreateThread(
        NULL, /* default security attributes */
        0, /* default stack size */
        Summation, /* thread function */
        &Param, /* parameter to thread function */
        0, /* default creation flags */
        &ThreadId); /* returns the thread identifier */

    if (ThreadHandle != NULL) {
        /* now wait for the thread to finish */
        WaitForSingleObject(ThreadHandle, INFINITE);

        /* close the thread handle */
        CloseHandle(ThreadHandle);

        printf("sum = %d\n", Sum);
    }
}
```

Java Threads

```
class Sum
{
    private int sum;

    public int getSum() {
        return sum;
    }

    public void setSum(int sum) {
        this.sum = sum;
    }
}

class Summation implements Runnable
{
    private int upper;
    private Sum sumValue;

    public Summation(int upper, Sum sumValue) {
        this.upper = upper;
        this.sumValue = sumValue;
    }

    public void run() {
        int sum = 0;
        for (int i = 0; i <= upper; i++)
            sum += i;
        sumValue.setSum(sum);
    }
}
```

```
public class Driver
{
    public static void main(String[] args) {
        if (args.length > 0) {
            if (Integer.parseInt(args[0]) < 0)
                System.err.println(args[0] + " must be >= 0.");
            else {
                Sum sumObject = new Sum();
                int upper = Integer.parseInt(args[0]);
                Thread thrd = new Thread(new Summation(upper, sumObject));
                thrd.start();
                try {
                    thrd.join();
                    System.out.println
                        ("The sum of "+upper+" is "+sumObject.getSum());
                } catch (InterruptedException ie) { }
            }
        }
        else
            System.err.println("Usage: Summation <integer value>");
    }
}
```

Threading Issues

- Semantics of `fork()` and `exec()` system calls
 - If one thread in a program calls `fork()`, does the new process duplicate all threads, or is the new process single-threaded
- Thread cancellation
 - The difficulty with cancellation occurs in situations where resources have been allocated to a canceled thread or where a thread is canceled while in the midst of updating data it is sharing with other threads
- Signal handling
 - a signal is typically delivered only to the first thread found that is not blocking it.
- Thread-local storage
- Scheduler activations